Field	Details
Product	Darwinbox — Enterprise HCM Platform
Feature Name	HireIQ — AI-Powered Hiring Intelligence
Primary User	HR Manager (Enterprise, 500–5,000 employees)
AI Type	LLM + RAG on internal HR + performance data
Document Status	Discovery Complete → Solution Framing → P0 Defined
Author	Abhishek Jha
Appetite	10–12 weeks build · P0 scope locked

1. Strategic Context

Darwinbox competes in India's enterprise HCM market against global players:

Competitor	Strength	Darwinbox Differentiation
SAP SuccessFactors	Global enterprise depth + compliance	India-first + modern UX + faster implementation
Workday	Best-in-class analytics + finance integration	Price + flexibility + regional adaptability
PeopleStrong	India domain knowledge	Superior UI + broader module coverage
Keka HR	SMB focus + ease of use	Enterprise depth + scalability
Oracle HCM	ERP integration + global scale	Agility + AI innovation velocity

While Darwinbox has built strong differentiation in core HCM modules, the talent acquisition layer remains largely transactional — it records hiring activity but does not meaningfully predict hiring outcomes.

> **The Strategic Opportunity:**
>
> Every enterprise HR manager has the same unsolved problem: they hire people based on resumes, interviews, and gut instinct — then spend 6–18 months discovering if the hire was right.
>
> Darwinbox already holds the answer inside its own platform: years of performance data, tenure records, manager effectiveness scores, and team composition data for thousands of employees across its enterprise clients.
>
> **HireIQ turns that historical data into predictive hiring intelligence.** Not 'does this candidate look good on paper?' — but 'will this specific candidate actually succeed in this specific role, under this specific manager, in this specific team?'
>
> This is the shift from **recording hiring** to **intelligently predicting hiring outcomes.**

2. Core Problem (Validated)

HR Managers at enterprises making 50–200 hires per year face a consistent, expensive, and largely invisible problem:

> **Bad hires are invisible until it is too late.**
>
> The average HR manager selects candidates based on resume signals, structured interview scores, and reference checks. These inputs are weak predictors of actual job performance — especially at 6, 12, and 18 month horizons.
>
> Critically, the data that WOULD predict success — how similar past employees performed in identical roles under the same managers — already exists inside Darwinbox. But it is never surfaced during the hiring decision.
>
> **This is not a data problem. It is a decision intelligence problem.**

Root Causes — The Five Friction Multipliers:

#	Root Cause	Impact	Frequency
1	No connection between historical performance data and current hiring decisions	Critical — directly causes repeat hiring mistakes	Every hire
2	Interview scores are subjective and interviewer-dependent	High — same candidate scores differently across interviewers	Every hire
3	No manager-fit signal — same candidate may thrive under one manager, fail under another	Critical — manager fit is #1 predictor of 1-yr retention	Every hire
4	Onboarding success patterns not fed back into hiring criteria	High — learning from past hires never reaches hiring stage	Every cycle
5	HR manager has no 'confidence score' — acts on gut when data exists	Medium — creates regret, re-hiring cost, team disruption	Most hires

The Cost of Getting This Wrong:

Hire Level	Replacement Cost	Time to Detect Bad Hire	Cascade Impact
Entry Level (0–2 yrs)	50% of annual salary	6–9 months	Team morale + manager time
Mid Level (3–6 yrs)	100% of annual salary	9–12 months	Project delays + client risk
Senior Level (7+ yrs)	200% of annual salary	12–18 months	Strategic initiative failure
Leadership	400% of annual salary	18–24 months	Culture + business continuity

For a 2,000 person enterprise making 300 hires/year with 20% mis-hire rate = 60 bad hires = 15–40 Crore annual cost.

3. Divergence — Full Solution Landscape

We explored 18 solution concepts across 5 problem dimensions before converging on the P0 scope:

Dimension	Solutions Explored	Selected?
Candidate Prediction	Success Score, Role Fit Index, Manager Fit Score, Team Fit Score, Culture Match	**YES — P0**
Conversational AI	Natural language candidate Q&A, AI interview debrief, HR chatbot for hiring	**YES — P0**
Decision Support	Confidence Score, Red Flag Alerts, Comparison Panel, Decision Nudge	**YES — P0**
Automation	Auto-JD generation, Resume auto-ranking, Interview scheduling AI	**P1**
Analytics	Hiring funnel AI, Source quality predictor, Offer acceptance predictor	**P1**

We deliberately included low-confidence and moonshot ideas to avoid premature convergence. The full exploration generated 18 concepts, and RICE-based prioritisation reduced the the scope to 3 P0 features.

4. Prioritisation Logic (RICE-Based)

RICE = (Reach x Impact x Confidence) / Effort. Scale: 1–3 for each dimension.

Feature	Reach	Impact	Confidence	Effort	RICE	Decision
Candidate Success Score	3	3	3	2	**13.5**	**P0**
Manager Fit Intelligence	3	3	2	2	**9.0**	**P0**
AI Hiring Assistant (RAG)	3	3	2	2	**9.0**	**P0**
Auto JD Generator	2	2	2	2	4.0	P1
Interview Scheduling AI	2	2	2	3	2.7	P1
Source Quality Predictor	2	2	1	2	2.0	P2
Offer Acceptance Predictor	2	2	1	3	1.3	P2
Culture Match Score	1	2	1	2	1.0	Deferred

Top 3 P0 Features directly attack:

- Manual, intuition-based candidate comparison (Success Score)
- Manager-fit blind spot — the #1 unaddressed predictor (Manager Fit Intelligence)
- Inability to query historical hiring data in real time (AI Hiring Assistant)

5. Final P0 Scope (Appetite: 10–12 Weeks)

We are committing to three tightly scoped features that together deliver a complete hiring intelligence layer:

P0.1 — Candidate Success Score

An AI-generated probability score (0–100%) predicting candidate success in a specific role, based on patterns from similar past hires in the organisation's own Darwinbox data.

Component	Detail
Score Range	0–100% success probability
Model Type	Supervised ML — trained on historical employee data (performance, tenure, attrition)
Input Signals	Role match, skills alignment, education pattern, past tenure, interview score, manager history
Output	Success Score + Top 3 contributing factors + Benchmark comparison (vs avg hire in this role)
Confidence Band	High / Medium / Low — shown alongside score to prevent over-reliance
Tooltip	"Based on 47 similar past hires at your company, candidates with this profile averaged 18-month tenure"
Fallback	If insufficient historical data: show industry benchmark with a clear label

P0.2 — Manager Fit Intelligence

An AI signal showing how well a candidate's work style aligns with the hiring manager's team patterns — based on performance and retention data of the manager's past direct reports.

Component	Detail
Signal Type	Manager Fit Score (High / Medium / Low) + narrative explanation
Data Source	Manager's historical team: retention rates, performance distributions, feedback themes
Output	"This manager's team retains 78% of hires at 12 months. Candidates with profile X perform best."
Privacy Design	Aggregated team-level data only. Individual employee records never surfaced.
Minimum Data Bar	Requires a minimum of 5 past direct reports to generate a signal. Below threshold: show N/A with explanation.
Bias Guardrail	Demographic attributes were excluded from all model features. Bias audit pre-launch required.

P0.3 — AI Hiring Assistant (LLM + RAG)

A conversational AI interface embedded in the hiring workflow that answers HR manager questions using the organisation's own Darwinbox data — without requiring manual report generation.

Component	Detail
Interface	Chat panel embedded in candidate profile view — collapsible, non-intrusive
Architecture	LLM + RAG: candidate + employee data chunked, embedded, stored in vector DB. Query → retrieve → generate.
Sample Queries	"How did our last 5 hires for this role perform at 6 months?" / "What profiles succeeded under this manager?" / "Are there red flags in this candidate's profile?"
Data Scope	Reads: candidate data, historical employee records (aggregated), performance data, manager profiles
Writes	None — read-only. All actions require explicit HR manager confirmation.
Citation Design	Every AI answer cites its source: "Based on 12 employee records in your system"
Confidence	Low confidence answers flagged explicitly: "Limited data available — treat this as directional only"
Escalation	Sensitive queries (compensation, legal) auto-escalate to HR note with disclaimer

6. User Persona

Priya Sharma — HR Manager, TechCorp India (1,800 employees)

Dimension	Detail
Age / Role	34 years old. HR Manager. Reports to CHRO. Manages a team of 3 HR executives.
Hiring Volume	Manages 80–120 hires per year across engineering, sales, and operations functions.
Daily Reality	Spends 60% of time on hiring coordination. Reviews 15–30 candidate profiles per week.
Current Process	Reviews resume in Darwinbox ATS → Coordinates 3-round interviews → Collects scores → Makes offer decision
Biggest Pain	"I make the best decision I can with what I have. But I never really know if I am making the right call until 12 months later."
Secondary Pain	"I have hired the same profile twice for the same manager and had completely different outcomes. I do not know why."
Tech Comfort	High. Uses Darwinbox daily. Comfortable with dashboards. Expects clean, non-intrusive AI features.
Trust in AI	Moderate. Wants AI as a second opinion, not a decision maker. Needs to understand WHY.
Success Definition	"If I can reduce my 20% bad-hire rate to 10%, that is a massive win for the business and for my credibility."

7. User Stories & Acceptance Criteria

US-01 — Candidate Success Score

User Story: As an HR Manager reviewing a candidate profile, I want to see an AI-generated success probability score so that I can make a more informed shortlisting decision.

Acceptance Criteria:

- Success Score (0–100%) visible on candidate profile card in ATS view
- Top 3 contributing factors displayed below score with plain English explanation
- Confidence band (High/Medium/Low) shown alongside score
- Benchmark comparison shown: 'vs avg hire in this role at your company'
- Tooltip on hover: data source and sample size displayed
- Score updates within 2 seconds of candidate profile load
- If insufficient data: fallback message displayed — not an empty state

US-02 — Manager Fit Intelligence

User Story: As an HR Manager, I want to see how well a candidate fits the hiring manager's team style so that I can flag potential mismatches before the offer stage.

Acceptance Criteria:

- Manager Fit Signal (High/Medium/Low) visible on candidate profile
- Narrative explanation: 'This manager retains X% of hires. Profile Y performs best.'
- Signal only shown when manager has minimum 5 historical direct reports
- Below minimum threshold: 'Insufficient data' shown — not hidden
- Demographic attributes never shown in signal explanation
- Bias audit passed before feature goes live — documented and accessible
- HR manager can click 'Why this score?' for expanded explanation

US-03 — AI Hiring Assistant — Basic Query

User Story: As an HR Manager, I want to ask natural language questions about this candidate and similar past hires so that I can quickly access relevant historical context without running manual reports.

Acceptance Criteria:

- Chat interface accessible from candidate profile — single click to open
- Responds to natural language queries in under 3 seconds
- Every answer cites data source with record count
- Low confidence answers explicitly flagged in UI
- Read-only — no data modifications possible via chat
- Sensitive query types (compensation, legal) trigger disclaimer + escalation note
- Chat history persists within session — clears on profile change

US-04 — AI Hiring Assistant — Red Flag Detection

User Story: As an HR Manager, I want the AI to proactively surface any red flags in a candidate's profile so that I do not miss patterns I might overlook manually.

Acceptance Criteria:

- AI proactively flags: tenure gaps, repeated short stints, role-skill mismatch
- Red flags shown as non-blocking alerts — informational, not gatekeeping
- Each flag includes plain English explanation and historical context
- HR manager can dismiss flag with one click — dismissal logged for model feedback
- False positive rate monitored post-launch — threshold: < 15% dismissal rate
- No flag generated for protected class attributes — bias review pre-launch

User Story: As an HR Manager, I want to be able to override AI recommendations and provide feedback on prediction quality so that the model improves over time and I retain decision authority.

Acceptance Criteria:

- Override button available on every AI-generated score
- Override requires optional reason selection (3 options + free text)
- Override logged for model retraining pipeline
- Post-hire outcome feedback requested at 3-month and 12-month marks
  Feedback loop visible to HR manager: 'Your feedback improved X predictions this quarter'
- Override rate monitored: if > 40% override rate — model review triggered
-

US-05 — Human Override + Feedback

8. AI Model Design & Architecture

8.1 — The HireIQ Technical Stack

Layer	Component	Technology	Purpose
Data Layer	Employee Records	Darwinbox HRIS	Performance, tenure, attrition history
Data Layer	Manager Profiles	Darwinbox People Mgmt	Team retention, performance distributions
Data Layer	ATS Data	Darwinbox ATS	Candidate profiles, interview scores
ML Layer	Success Score Model	XGBoost / LightGBM	Supervised classification — predict hire success
ML Layer	Feature Store	Feast / Custom	Pre-computed candidate + role features
RAG Layer	Embedding Model	OpenAI Ada / BGE	Convert HR records to vectors
RAG Layer	Vector Database	Pinecone / pgvector	Store + retrieve relevant HR context
LLM Layer	Generation Model	GPT-4 / Claude	Generate natural language answers
Product Layer	API Gateway	REST API	Connect models to Darwinbox UI
Product Layer	Explainability	SHAP Values	Generate factor explanations for scores

8.2 — The RAG Flow for AI Hiring Assistant

Step 1 — Ingestion (runs nightly): Employee records, performance data, and manager profiles are chunked into semantic units and embedded using a text embedding model. Stored in vector database with metadata (role, department, manager, date).

Step 2 — Query Processing: HR Manager types a question. Query is embedded using same model. Top-k most semantically relevant chunks retrieved from vector DB. Relevance threshold: cosine similarity > 0.75.

Step 3 — Generation: Retrieved context + original query sent to LLM with system prompt defining: role (HR intelligence assistant), constraints (read-only, cite sources, flag uncertainty), and tone (professional, concise, non-alarmist).

Step 4 — Response Display: Answer shown with source citation (record count, data type, date range). Confidence level displayed. Sensitive topic triggers disclaimer. All responses logged for quality review.

8.3 — Success Score Model Design

Design Decision	Choice Made	Rationale
Model Type	Supervised Binary Classification	Predict: Will this hire succeed in 12 months? (Yes/No threshold)
Training Label	12-month performance + retention	Most reliable signal of hire quality available in the Darwinbox data
Key Features	Role match, skills overlap, past tenure, interview score, manager fit score, team composition	All available in Darwinbox — no external data required
Excluded Features	Gender, age, education institution, location, name	Bias prevention — demographic proxies removed
Minimum Data	50 past hires per role family minimum	Below threshold: fallback to industry benchmark with clear label
Evaluation Metric	Recall > 70% (primary), Precision > 60% (secondary)	Missing a bad hire (low recall) is more costly than a false flag
Bias Audit	Demographic parity across gender, age, tenure groups	Required before launch — must pass before feature goes live
Retraining	Monthly — using outcome feedback from 3-month + 12-month surveys	Keeps model fresh as hiring patterns evolve

9. Edge Cases & UX Design

Edge Case	Trigger	UX Response	Why This Matters
Insufficient Data	< 50 past hires for this role	Show industry benchmark with clear label: 'No internal data yet — showing industry pattern'	Prevents empty state — always gives HR manager something actionable
Low Model Confidence	Confidence band = Low	Show score with prominent Low Confidence badge + 'Use as one signal among many'	Prevents over-reliance on uncertain predictions
New Manager	Manager has < 5 historical direct reports	Manager Fit signal hidden — replaced with 'Building profile — check back after first 5 hires'	Prevents showing unreliable signal as credible data
AI Assistant Hallucination	LLM generates an unsupported claim	System prompt + RAG grounding + confidence check before display. Flagged answers shown with uncertainty label	Trust is everything — one wrong answer destroys feature adoption
Sensitive Query	Compensation, legal, discrimination keywords	Auto-disclaimer shown: 'This may involve sensitive data. Please verify with HR policy before acting.'	Protects the HR manager from compliance risk
Model Unavailable	API timeout or model error	Graceful degradation: hide AI features, show base ATS view, silent error logging	The core ATS function is never blocked by an AI layer failure
Override Rate Spike	Override rate > 40% in 30 days	Automated alert to PM + model review triggered within 5 business days	Early warning system for model drift or trust erosion
Bias Alert	Demographic parity score drops below threshold	Feature paused automatically + PM + legal notified within 24 hours	Zero tolerance for biased outcomes in hiring decisions

10. What We Explicitly Did NOT Include (Yet)

Excluded Feature	Reason for Exclusion	Future Potential
Auto-hire decision making	Legal + ethical — AI must advise, never decide on people outcomes	Never — human in loop is permanent
Multi-platform data import	Data trust risk + integration complexity exceeds P0 appetite	P2 — after internal data quality proven
Video interview AI analysis	High regulatory risk (EU AI Act High Risk) + accuracy concerns	Deferred — pending regulation clarity
Salary prediction	Compensation decisions require separate governance + legal review	Separate initiative with legal sign-off
External LinkedIn data	Privacy risk + consent complexity + data freshness issues	P2 — with privacy-first design
Candidate-facing AI	Candidate trust + bias transparency requires separate design sprint	P1 — after HR-facing proven
Automated interview scheduling	Lower RICE score + lower strategic differentiation	P1 — operational efficiency feature

Clear boundaries prevent scope creep and protect delivery integrity

11. User Journey (Integrated Flow)

Step	User Action	System Response	AI Involvement
1	HR Manager opens candidate profile in Darwinbox ATS	Standard candidate profile loads with all existing ATS data	None — base experience unchanged
2	Candidate Success Score loads automatically	Score (0–100%) appears with confidence band + top 3 factors + benchmark comparison	ML model — Success Score
3	Manager Fit Intelligence loads	Manager Fit signal (H/M/L) appears with narrative explanation	ML model — Manager Fit Score
4	HR Manager wants deeper context	Opens AI Hiring Assistant chat panel — asks: 'How did similar hires perform for this manager?'	LLM + RAG — retrieves and generates answer with citation
5	AI surfaces a red flag proactively	Non-blocking alert: '3 similar profiles showed short tenure in operations roles'	ML anomaly detection — pattern matching
6	HR Manager reviews all signals	Has: Success Score + Manager Fit + AI answer + Red Flag context	All three P0 features active simultaneously
7	HR Manager makes decision	Proceeds to shortlist / reject / hold — decision is 100% HR manager's	AI advises. Human decides. Always.
8	Outcome feedback at 3 months	The system sends the HR manager: 'How is this hire performing?'	Feedback logged → model retraining pipeline

No forced paths. All AI features are additive and non-blocking. Base ATS experience is preserved at all times

12. Metrics Framework (Expanded)

North Star Metric:

12-Month Hire Success Rate

Percentage of AI-assisted hires that are still employed and rated 'meeting expectations' or above at 12 months.

Target: +15% improvement over baseline within 6 months of full rollout.

Why this metric: It is the only metric that connects AI accuracy to real business value — reducing replacement cost and improving team stability.

Secondary Metrics:

Metric	Target	Why It Matters	Measured How
Success Score CTR	> 70% of HR managers view score	Adoption signal — is the feature being used?	Darwinbox event tracking
AI Assistant Query Rate	> 3 queries per hiring decision	Engagement signal — is the RAG useful?	Chat session analytics
Override Rate	< 30% of scores overridden	Trust signal — do HR managers trust the predictions?	Override logging
Time to Shortlist Decision	Decrease by 20%	Efficiency signal — does AI reduce decision friction?	ATS workflow timing
Manager Fit Signal Usage	> 50% view rate	Adoption of a differentiated feature	Event tracking
Red Flag Dismissal Rate	< 15% dismissal rate	Quality signal — are alerts relevant and actionable?	Dismissal logging
Feedback Completion Rate	> 60% at 3-month mark	Data quality signal for model retraining	Survey completion tracking

Guardrail Metrics:

Guardrail	Threshold	Trigger Action
Demographic Parity Score	Must stay within 5% across gender, age, tenure groups	Feature paused + PM + legal notified within 24 hours
Override Rate Spike	If > 40% override rate in any 30 days	Model review triggered within 5 business days
AI Answer Accuracy	If > 20% of RAG answers are flagged as inaccurate by users	RAG pipeline review + prompt refinement sprint
12-Month Cancellation Rate	AI-assisted hires must not have a higher early resignation rate	Full model audit if regression is detected
Model Latency	Score must load in < 2 seconds P95	Engineering alert + performance sprint

13. Experimentation Plan

Phase	Scope	Features Live	Success Threshold	Duration
Phase 1 — Shadow Mode	Internal only — no HR manager visibility	All 3 features run silently alongside the current process	Model accuracy > 70% Recall vs actual outcomes	8 weeks
Phase 2 — Pilot	3 design partner clients (500–2,000 employees)	Success Score + AI Assistant only	NPS > 40, Override rate < 35%, No bias flags	4 weeks
Phase 3 — Limited Roll	25% of Darwinbox enterprise clients	All 3 P0 features	> 5% improvement in 3-month hire satisfaction	6 weeks
Phase 4 — Full Launch	All enterprise clients	All 3 P0 features + feedback loop active	North Star: +15% 12-month hire success rate	Ongoing

Rollback Capability at Every Phase:

- Success Score: disable display without affecting ATS core function
- Manager Fit Signal: toggle off per client — no system dependency
- AI Assistant: feature flag off — chat panel hidden, zero ATS impact
- Each feature independently deployable and independently rollback-able

14. Risk Matrix

Risk	Probability	Impact	Mitigation	Owner
HR managers don't trust AI scores	High	High	Explainability first — always show WHY. Confidence bands. Human override. Pilot with power users first.	PM + Design
Model perpetuates hiring bias	Medium	Critical	Demographic exclusion in features. Pre-launch bias audit. Demographic parity guardrail. Legal review.	PM + Legal + Data Science
Insufficient historical data for new clients	High	Medium	Industry benchmark fallback with clear label. Minimum data threshold before score shown.	PM + Engineering
RAG hallucination damages trust	Medium	High	System prompt constraints. RAG grounding. Confidence scoring. Source citation mandatory.	PM + Data Science
Privacy concerns about employee data usage	Medium	High	Aggregated data only in Manager Fit. Clear data usage disclosure. DPDP Act compliance.	PM + Legal
Slow model performance kills UX	Low	Medium	Pre-computed scores (nightly batch). < 2s load time SLA. Engineering sign-off pre-launch.	Engineering
Scope creep from stakeholders	High	Medium	P0 scope locked in this document. Change requests require RICE re-evaluation.	PM
EU AI Act high-risk classification	Medium	High	Document all model decisions. Human oversight mandatory. Conformity assessment prep.	PM + Legal

15. Technical Considerations

Area	Requirement	Engineering Sign-off Required
Data Pipeline	Nightly batch job: employee records → feature engineering → model inference → score storage	Yes — before Phase 1
Vector Database	Employee + candidate data embedded and stored. Query latency < 500ms P95.	Yes — before Phase 2
LLM Integration	API integration with chosen LLM provider. System prompt design + prompt injection prevention.	Yes — before Phase 2
SHAP Integration	SHAP values computed per prediction. Top 3 factors extracted and formatted for display.	Yes — before Phase 2
Feature Flags	Per-client, per-feature toggle capability. All 3 P0 features independently controllable.	Yes — before Phase 3
Bias Monitoring	Real-time demographic parity monitoring. Automated alert + feature pause if threshold breached.	Yes — before Phase 3
Caching Layer	Pre-computed scores cached per candidate per role. Cache invalidated on profile update.	Yes — before Phase 3
DPDP Compliance	Data processing consent framework. Employee right to explanation. Data retention policy.	Yes — before full launch

16. Decision Log

Decision	What We Chose	What We Rejected	Rationale
Score display format	0–100% with confidence band	Green/Amber/Red traffic light only	Numeric gives HR manager more calibration. Confidence band prevents over-reliance.
Data scope for Manager Fit	Aggregated team data only	Individual employee performance data	Privacy protection. Aggregation maintains signal without exposing individuals.
RAG vs fine-tuning	RAG architecture	Fine-tuned domain model	HR data changes constantly. RAG stays current. Fine-tuning would be stale within months.
AI decides vs AI advises	AI advises only — HR manager decides	Automated shortlisting	Legal + ethical. Employment decisions require human accountability.
Shadow mode first	8 weeks shadow before HR visibility	Direct pilot launch	Trust in model quality before exposing to users. Prevents early bad experience.
3 P0 features vs 1	3 interconnected features	Single feature MVP	Each feature alone is weak. Together: Score + Fit + Assistant = complete intelligence layer.

17. Rollout Strategy

Stage	Scope	Criteria to Advance	Timeline
Stage 0 — Shadow Mode	Internal Darwinbox data only	Model Recall > 70%. No bias flags.	Weeks 1–8
Stage 1 — Design Partners	3 enterprise clients (opt-in). HR managers only.	NPS > 40. Override rate < 35%. No support spike.	Weeks 9–12
Stage 2 — Controlled Roll	25% of enterprise clients. HR managers only.	North Star moving +5%. No guardrail breaches.	Weeks 13–18
Stage 3 — Full Launch	All enterprise clients.	North Star: +15% hire success rate confirmed.	Week 19+
Stage 4 — P1 Expansion	Add candidate-facing features + JD AI.	P0 metrics stable for 90 days.	Quarter 3

18. Post-Launch Learning Agenda

We will evaluate the following questions in the 90 days post full launch:

Question	Data Source	Action If Negative
Are Success Score predictions aligned with actual 3-month hire quality?	3-month feedback survey	Model retraining + feature weight review
Is Manager Fit signal reducing manager-candidate mismatch?	12-month retention by manager	Manager Fit model architecture review
Are HR managers using AI Assistant or ignoring it?	Chat session analytics	UX redesign + onboarding improvement
Is bias audit holding at scale?	Demographic parity monitoring	Feature pause + model audit
Are bad-hire rates actually decreasing?	12-month performance outcomes	Full P0 review if no movement
Are there query types the AI Assistant cannot handle well?	Low confidence rate + dismissals	RAG pipeline expansion + prompt refinement

Next Potential Iterations (P1 Roadmap):

Feature	Trigger Condition	Value
Candidate-facing success insights	P0 HR trust established. Bias audit passed.	Candidate transparency + Darwinbox differentiation
JD AI Generator	P0 adoption > 60% of HR managers	Reduce time-to-JD by 70%
30-day onboarding intelligence	3-month feedback loop proven	Early intervention before bad hire is confirmed
Cross-client benchmarking	5+ clients on P0 with outcome data	Industry benchmarks for roles without internal data
Interview question AI	Manager Fit signal trusted by > 70% of managers	Personalised interview guides per candidate

19. Final Summary

HireIQ reframes Darwinbox from:

"A system that records who you hired"

→

"A system that tells you who will actually succeed before you hire them."

By turning Darwinbox's own historical data into a predictive intelligence layer, HireIQ attacks the #1 most expensive, most invisible problem in enterprise HR: the bad hire that nobody saw coming.

What HireIQ Delivers	Business Outcome
Candidate Success Score	Reduce mis-hire rate by up to 50% using historical pattern matching
Manager Fit Intelligence	Address #1 unaddressed predictor of 12-month retention
AI Hiring Assistant (RAG)	Zero manual report generation — HR context in seconds
Bias Audit + Human Override	Responsible AI that HR teams trust and regulators accept

This is not an AI feature bolted onto a hiring tool.

This is a fundamental upgrade to how enterprise hiring decisions are made.